

## The Super Loop

```
int main(void) {  
    /* Initialize system */  
  
    while (1) {  
        /* Periodic Tasks */  
        ADC_Read();  
        SPI_Read();  
        USB_Packet();  
        Audio_Decode();  
        File_Write();  
    }  
}
```


“No interrupts”  
“Polling”

+ Simple  
Minimal HW resources  
Highly portable  
  
- Inaccurate timing  
High power consumption

- 1 - Determinism
- 2 - Responsiveness
- 3 - Polling periodicity

## Foreground/Background

```
void USB_ISR(void) {  
    Clear interrupt;  
    Read packetM  
}
```



- + No upfront cost  
Minimal training required  
No need to set aside resources to accommodate RTOS
- Difficult to ensure that each operation will meet its deadline  
High-priority code must be placed in the foreground  
Problems can arise when code is maintained by multiple developers  
Even in projects with a single developer, expanding the application can prove difficult

## ISR

- An ISR refers to the Interrupt Service Routine. These are procedures stored at specific memory addresses which are called when a certain type of interrupt occurs
- An ISR returns nothing and not allowed to pass any parameters. An ISR is called a hardware or software event occurs, it is not called by the code, so that's the reason no parameters are passed into an ISR
- Interrupt latency is the number of clock cycles that is taken by the processor to respond to an interrupt request. These number of cycles is counted between the assertions of the interrupt request and first instruction of the interrupt handler

- You can measure interrupt latency with the help of an oscilloscope. Take two GPIOs - one for interrupt and second for the toggling
- Reduce the interrupt latency
  - Platform and interrupt controller
  - CPU clock speed
  - Timer frequency
  - Cache configuration
  - Application program
- Causes of interrupt latency
  - The interrupt request signal needs to be synchronized to the CPU - upto 3 CPU cycles before the interrupt has reached the CPU core
  - The CPU will typically complete the current instruction, which may take several cycles
  - After completion of the current instruction, the CPU performs a mode switch or pushes registers on the stack
  - Pipeline fill. An instruction is executed when it has reached its final stage of the pipeline. Since the mode switch has flushed the pipeline, a few extra cycles are required to refill the pipeline
- Functions which are called from the ISR should be reentrant

## **RTOS**

RT = Correct function @ Correct Time

OS = HW + SW manager

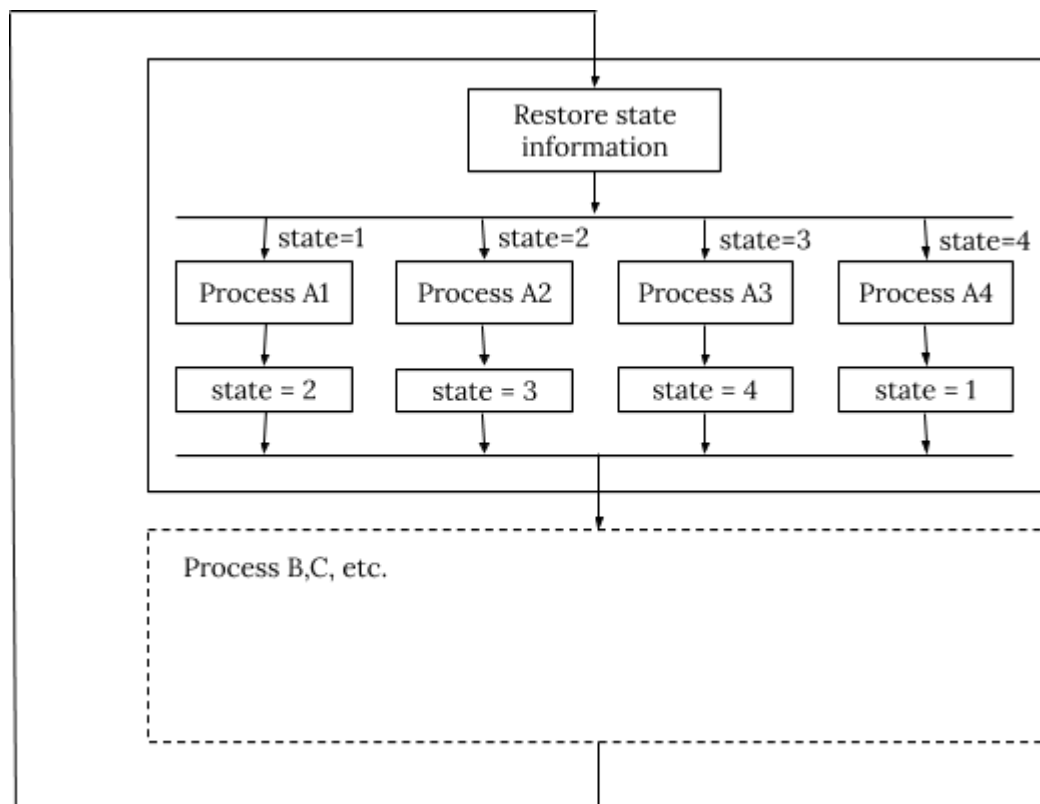
RTOS = HW + SW manager that can help us ensure having correct function @ correct time

Real-Time System

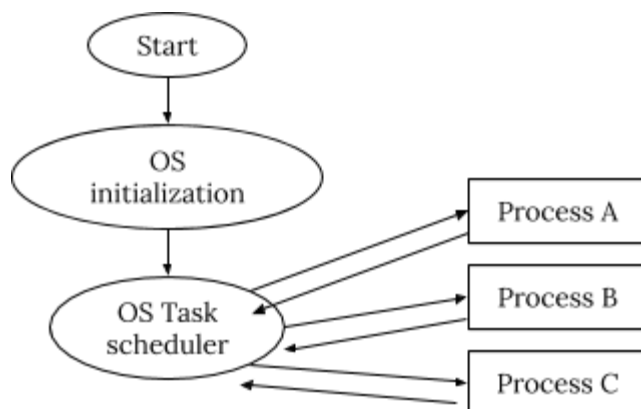
They can be:

- Hard
  - Firm
  - Soft
- Fast software is not necessarily real-time software
  - Determinism is a desirable quality in real-time software
  - Software that is deterministic has a bounded response time to events


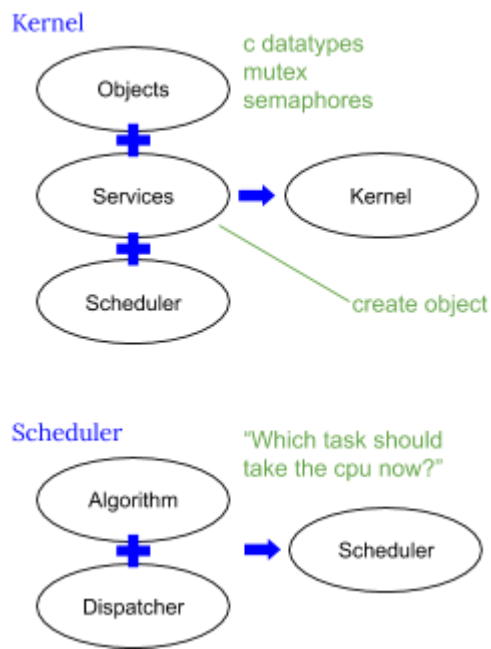
## **Concurrent Processes**



### Concurrent Processes: RTOS



**RTOS = Kernel + ...**



### Benefits


Developers who use RTOS are freed from implementing a scheduler and related services.

Typical applications that incorporate RTOS are much easier to expand.

The best RTOS have undergone thorough testing

### Types of Kernels

- Cooperative Kernel
- Preemptive Kernel



### To RTOS or not to RTOS

Technical and Managerial decision based on many factors including:

- Expected size of the code

- Complexity of the hardware
- Timing requirements
- Size of the development team

Foreground/Background systems that include any of the following should consider migrating to an RTOS

- Excessive Polling
- Counters for controlling execution rates
- Repetitive function calls in the main loop

```
void ADC_Read(void) {
    while((ADC_ConvComplete()) == 0) {
        ;
    }
    Process analog value;
}
```

Excessive Polling

```
while (1) {
    ADC_Read();
    if((i % 8192) == 0) {
        SPI_Read();
    }
    i++;
}
```

Counters

```
while(1) {
    ADC_Read();
    LCD_Update();
    SPI_Read();
    USB_Packet();
    LCD_Update();
    Audio_Decode();
    File_Write();
    LCD_Update();
}
```

Repetitive Function calls

### main()

- First function executed in  $\mu$ C/OS-II-based application
- The routines that main() calls lay the ground for multitasking

```
int main(void) {
    OS_Init();
    /* Create @ least 1 task */
    OS_Start();
}
```

### OSInit()

- OSInit() must be invoked before any of  $\mu$ C/OS-II's services are used
- This function initialize the operating system's data structures
- $\mu$ C/OS-II's internal tasks are created by OSInit()
- The extent of initializations depends on configuration constants

### $\mu$ C/OS-II's internal Tasks

#### Idle Task

- Runs when other tasks are unable to do so
- Invokes a hookfunction OSTaskIdleHook()
- Is automatically assigned the lowest possible priority

#### Statistics Task

- Runs periodically
- Records CPU utilization
- Is capable of checking the size of stacks
- Has a slightly higher priority than the Idle task

### Creating the First Task

- After OSInit() has been invoked, tasks can be created
- Most applications create just one task in main()
  - If multiple tasks are created, the statistics task will not function properly
- Additional tasks can be created at essentially anytime

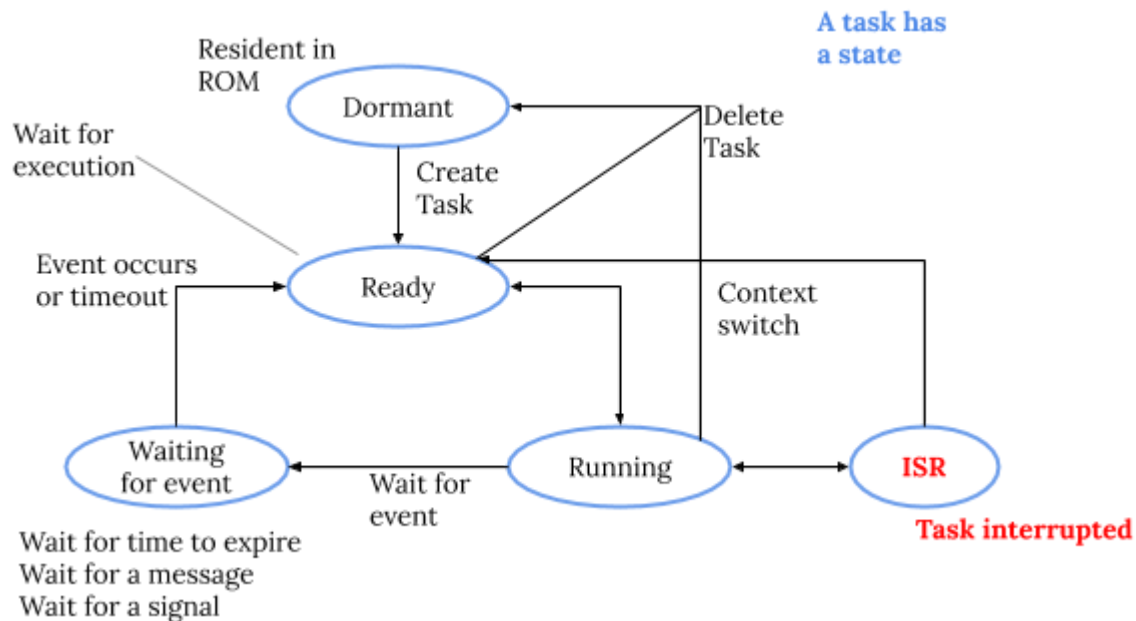
### OSStart()

- main() must call OSStart() in order to initiate multitasking
- OSStartHighRdy() which is an assembly-language routine invoked by OSStart() runs the first task
- Unless errors occurs, OSStart() should be the last function called from main()

### Task

- There are few required components
  - Each task is given its own stack space
  - A task must be assigned a priority
- Normally, a task involves an infinite loop
  - Task cannot return
- Initializations might precede the loop

```
void App_TaskExample(void *p_data) {
    Task initialiation;
    for(;;) {
        Work toward task's goal;
        Wait for events;
    }
}
```



### Task stack

- Each task should have its own stack
  - It must be declared as OS\_STK and made of consistent and contiguous memory
  - Can be allocated statically or dynamically
    - OS\_STK MyStack[??]

### Stack sizing


- Not all tasks requires the same amount of stack space
- Application developers are responsible for sizing their stacks

### Task Control Block


- TCB's for task management
- TCB's reside in RAM
- Every task is assigned a TCB when created
  - The task's priority
  - The state of the task
  - A pointer to the task's stack

### Application Partitioning

- Non-trivial
- A poorly partitioned application may fail to meet performance requirements
- Look for activities that can execute in parallel
- Beware of excessive communication and large number of tasks
- Many methods exists



**VS**



### Divide and Conquer (appl. partitioning)

#### Divide

- Identify code based on the following:
  - IO bound
  - CPU bound
  - Periodicity
- Functions can be either IO or CPU bounded
- Functions can be periodic or aperiodic

#### Conquer (Groups)

- Group functions into tasks based on the following:
  - Function cohesion
  - Time cohesion
  - Periodic cohesion
- Function cohesion → single task or sequential tasks
- Time cohesion → Separate parallel tasks
- Periodic cohesion → Same task with counters or different tasks

### Application Partitioning Sanity Check

For each task regardless of the execution schedule

$$T \leq D \leq P$$

T = WCET (worst case execution time)

D = Deadline

P = Period

### Assigning Priorities

#### Gomaa Criterion

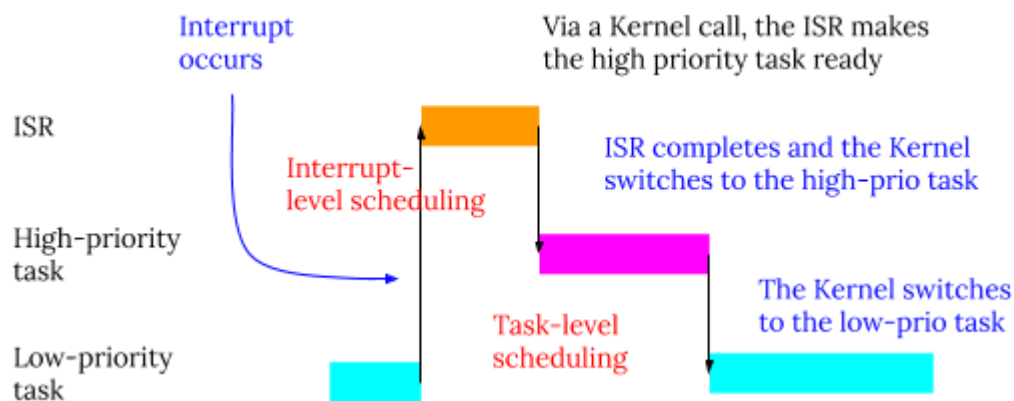
finish quickly    finish correctly



Urgent	Critical	Priority
No	No	Lowest
No	Yes	Lower
Yes	No	Higher
Yes	Yes	Highest



## RTOS scheduling



## Delay Functions

```
void OSTimeDly          (INT16U    ticks);
INT8U OSTimeDlyHMSM     (INT8U     hours,
                        INT8U     minutes,
                        INT8U     seconds,
                        INT16U     milli);
```

## Shared Resources

- A global variable or data structure that is used by multiple contexts
- Peripheral devices are shared resources
- Race condition, major problem in multi-context execution

## Protecting Shared Resources

- Shared resources are not atomically accessed and will never be atomically accessed
- Shared resources should be exclusively accessed
- Locking is the used mechanism
- Every shared resource is associated with a lock
  - A single lock can be used with multiple shared resources
- Access is granted to a context if it can open the lock

## Disabling Interrupts

- 2 macros provided OS\_ENTER\_CRITICAL() and OS\_EXIT\_CRITICAL()
- Used by  $\mu$ C/OS-II code
- Can be used by your application
  - Not a good coding style
  - Application may crash
- Can be implemented in 1 of 3 ways and selected through configuration

## Disabling Interrupts Method 3

- Compiler must support
  - Storing PSW in a C variable
  - Loading PSW from a C variable
  - Disabling interrupts in from C

```
void App_TaskExample(void* p_arg) {
#ifdef OS_CRITICAL_METHOD==3
```


```

    OS_CPU_SR cpu_sr = 0;
#endif
    while(1) {
        OS_ENTER_CRITICAL();
        Access shared resource;
        OS_EXIT_CRITICAL();
    }
}

```

### Schedule Locking

- A task can lock the scheduler to keep control of the cpu, even if there are higher priority tasks ready
- During scheduler locking, your application should not call any service that suspends execution
  - Application will crash!



### Semaphores

- Semaphores are based on counters
- A semaphore can be classified as either binary or counting
- Have 2 operations:
  - Pend = while the semaphore's counter has a value of zero, allow other tasks to run
  - Post = Increment the semaphore's counter
- Semaphores are implemented with event control blocks (ECBs)
  - These structures are somewhat similar to TCBs
- ECBs are used to implement semaphores, mutexes, mailboxes, and queues

### Semaphore API

```

OS_EVENT  *OSSemCreate    (INT16U  cnt);
void      OSSemPend       (OS_EVENT *pevent,
                           INT16U  timeout,
                           INT8U   *perr);
INT8U     OSSemPost       (OS_EVENT *pevent);

```

### Semaphore Problems

- 1 - Starvation
- 2 - Dead lock ⇒ Ordered locking
- 3 - Priority inversion

### Deadlock problem

```

Task1() {
Task2() {

```

```

lock 1
lock 2
unlock 2
unlock 1
}


```

```

lock 2 -> 1
lock 1 -> 2
unlock 1 -> 2
unlock 2 -> 1
}


```

## Priority Inversion



Mutex = Binary Semaphore used for Shared Resources Protection  
 +  
 Priority Inversion Solution  
 (Inheritance or ceiling)

## Priority Inheritance



## Mutex

- A mutex is yet another mechanism for protecting resources

- In  $\mu\text{C}/\text{OS-II}$ , mutexes provide built-in protection from priority inversion
  - A modified priority inheritance protocol is used
- Unlike a semaphore, a  $\mu\text{C}/\text{OS-II}$  mutex does not incorporate a counter
  - The mutex is either available or in use

```

OS_EVENT  *OSMutexCreate  (INT8U    prio,
                           INT8U    *perr);

void       OSMutexPend    (OS_EVENT  *pevent,
                           INT16U    timeout,
                           INT8U    *perr);

INT8U      OSMutexPost    (OS_EVENT  *pevent);

```

### Task Interaction

- The tasks comprising a  $\mu\text{C}/\text{OS-II}$ -based application are not necessarily self-contained
- In order for the application's objectives to be met, tasks may need to interact with each other (and possibly with ISRs)
- A typical RTOS provides services that facilitate such interaction

### Synchronizing a Task to an ISR

- Most applications must manage a collection of peripheral devices
- The interrupt service routines (ISRs), associated with system's peripheral devices should be kept brief
- In applications that incorporate a real-time kernel, ISRs can use synchronization primitives to signal tasks
- Using a semaphore, a task can synchronize to another task or to an ISR!

```

OS_EVENT *App_SemADC;
// Init code ...
App_SemADC = OSSemCreate(0);

```

```

void App_ISRADC(void) {
    Clear interrupt;
    OSSemPost(App_SemADC);
}

```

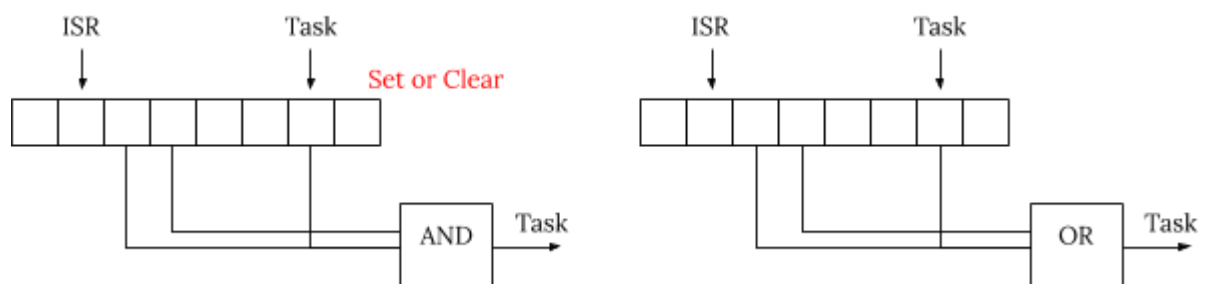
```

void App_TaskADC(void *arg) {
    perform init;
    while(1) {
        Start conversion;
        OSSemPend(App_SemADC, 0,
                  &err);
        Process converted value;
    }
}

```

### Event flags

- Using Event flags, a task can easily wait for multiple events to take place



```

OS_FLAG_GRP *OSFlagCreate    (OS_FLAGS      flags,
                               INT8U          *perr);
OS_FLAGS     OSFlagPend      (OS_FLAG_GRP    *pgrp,
                               OS_FLAGS      flags,
                               INT8U          wait_type,
                               INT16U        timeout,
                               INT8U          *perr);
OS_FLAGS     OSFlagPost      (OS_FLAG_GRP    *pgrp,
                               OS_FLAGS      flags,
                               INT8U          opt,
                               INT8U          *perr);

```

### Inter-Task Communication Services

- Tasks in a  $\mu$ C/OS-II-based application can send and receive messages using services that the kernel provides
- Application developers determine the content of these messages

### Message Queues

- In  $\mu$ C/OS-II, message queues are circular buffers
  - The kernel manages each buffer
  - Through API functions, tasks can request the insertion or removal of messages
- A message is a void pointer
- When a task is waiting on a message, the kernel runs other tasks

```

OS_EVENT     *OSQCreate      (void          **start,
                               INT16U        size);
void          *OSQPend       (OS_EVENT      *pevent,
                               INT16U        timeout,
                               INT8U          *perr);
INT8U         OSQPost        (OS_EVENT      *pevent,
                               void          *pmsg);

```

Q = Semaphore + circular buffer (Messages)

```

OS_EVENT *App_QUSB;
void      *App_BufUSB[APP_BUF_SIZE];
// init code
App_QUSB = OSQCreate(&AppBufUSB[0],
                    APP_BUF_SIE);


void App_ISRUSB(void) {
    Clear USB (or DMA) interrupt;
    OSQPost(App_QUSB, p_buf);
}

void App_TaskUSB(void *p_arg) {
    while(1) {
        p_buf = OSQPend(App_QUSB, 0, &err);
        process packet;
    }
}

```

}

### Many-to-1 communication



### Message Mailbox

- A message mailbox is another mechanism for inter-task communication
- Essentially, a mailbox is a queue that has a capacity of just one message
- In terms of overhead, mailboxes have a slight advantages over queues

OS_EVENT	*OSMboxCreate	(void	*pmsg);
void	*OSMboxPend	(OS_EVENT	*pevent,
		INT16U	timeout,
		INT8U	*perr);
INT8U	OSMboxPost	(OS_EVENT	*pevent,
		void	*pmsg);

### 1-Way Non-Interlocked Communication

```
OS_EVENT *App_MboxRPM;
// init code
App_MboxRPM = OSMboxCreate((void*) 0);

void App_ISRTimer(void) {
    Clear interrupt;
    Read timer value;
    OSMboxPost(App_MboxRPM, (void*)timer_val);
}

void App_TaskRPM(void *p_arg) {
    Initialize timer;
    while(1) {
        timer_val = (int *)OSMboxPend(App_MboxRPM, 0, &err);
        calculate RPM;
    }
}
```

### Memory Manager

- Fixed size memory block management
  - To prevent fragmentation
- Multiple partitions can be created with different sizes

- Blocks allocated from a certain partition must be returned back to the same partition

```

OS_MEM      *OSMemCreate      (void      *addr,
                                INT32U     nblks,
                                INT32U     blksize,
                                INT8U      *perr);

void         *OSMemGet         (OS_MEM     *pmem,
                                INT8U      *perr);

INT8U        OSMemPut          (OS_MEM     *pmem,
                                void        *pblk);

```

### Timer Management API

```

OS_TMR      *OSTmrCreate      (INT32U     dly,
                                INT32U     period,
                                INT8U      opt,
                                OS_TMR_CALLBACK callback,
                                void        *callback_arg,
                                INT8U      *pname;
                                INT8U      *perr);

BOOLEAN      OSTmrStart       (OS_TMR     *ptmr,
                                INT8U      *perr);

BOOLEAN      OSTmrStop        (OS_TMR     *ptmr,
                                INT8U      opt,
                                void        *callback_arg,
                                INT8U      *perr);

INT32U        OSTmrRemainGet   (OS_TMR     *ptmr,
                                INT8U      *perr);

INT8U         OSTmrStateGet    (OS_TMR     *ptmr,
                                INT8U      *perr);

```

### Endianness

- The Endianness is the order of bytes to store data in memory. If the system is big-endian then the MSB byte is stored first - little endian then LSB byte stored first

0x11223344

MSB      LSB

Big-endian

Address	Value
00	0x11
01	0x22
02	0x33
03	0x44

Little-endian

Address	Value
00	0x44
01	0x33
02	0x22
03	0x11

### RISC vs CISC

RISC = Reduced Instruction Set Computer

RISC processors have simple instructions taking about one clock cycle. The average clock cycle per instruction (CPI) is 1.5

Performance is optimized with more focus on software

RISC processors are highly pipelined

Code expansion can be a problem

The decoding of instructions is simple

RISC requires more RAM

CISC = Complex Instruction Set Computer

CISC processor has complex instructions that take up multiple clocks for execution. The average clock cycle per instruction (CPI) is in the range of 2 and 15.

Performance is optimized with more focus on hardware.

They are normally not pipelined or less pipelined

Code expansion is not a problem

Decoding of instructions is complex

CISC requires minimum amount of RAM